

OpenCL Tutorial



David Castells-Rufas


Microelectronics & Electronics Systems Department

Universitat Autònoma de Barcelona


david.castells@uab.cat

OpenCL for FPGAs

- We did



- But we need

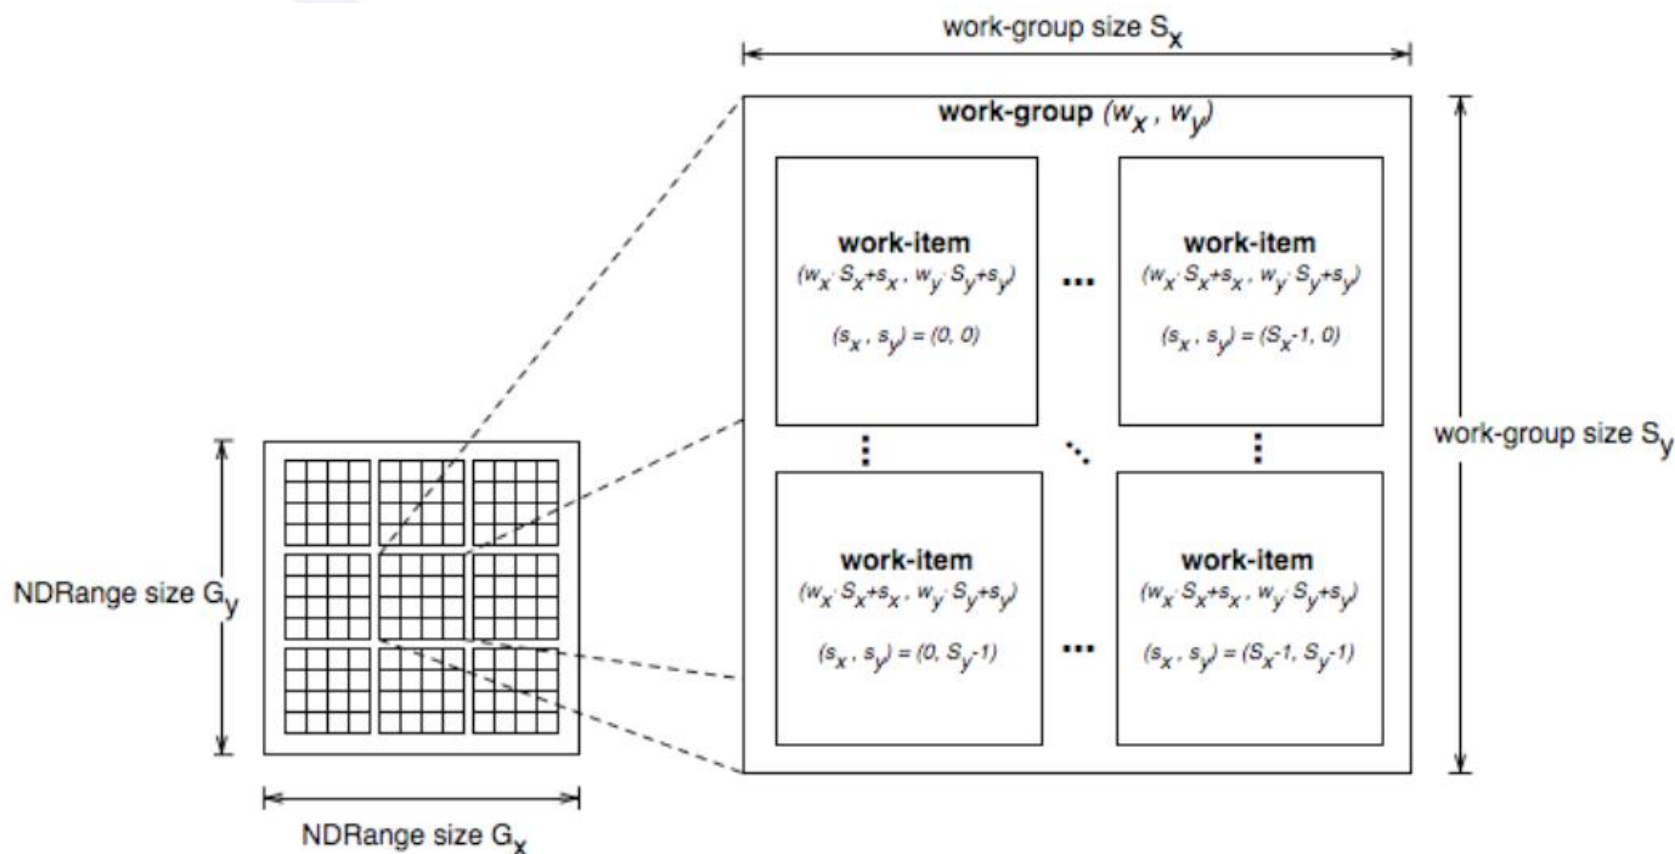


OpenCL

- API for the execution of “parallel” code (**kernels**) to be run in accelerators
- Programmers take care of parallelism
- OpenCL runtime
 - Compiles the code targeting the accelerator platform
 - Programs and executes (communicates with) the kernels in the accelerator platform

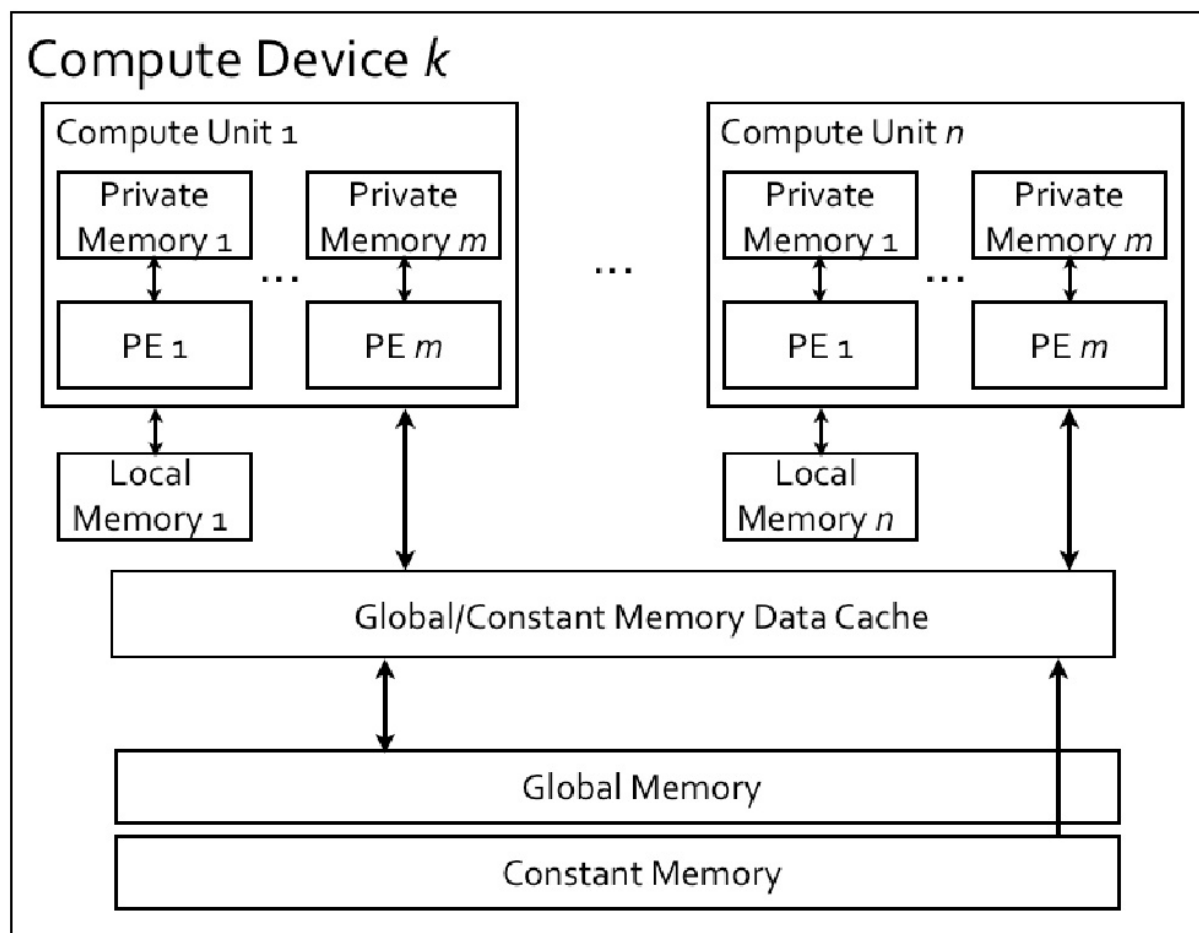
Logical organization of work

- Work Items (like threads)
- Work Groups (groups of threads executed in the same computing unit)



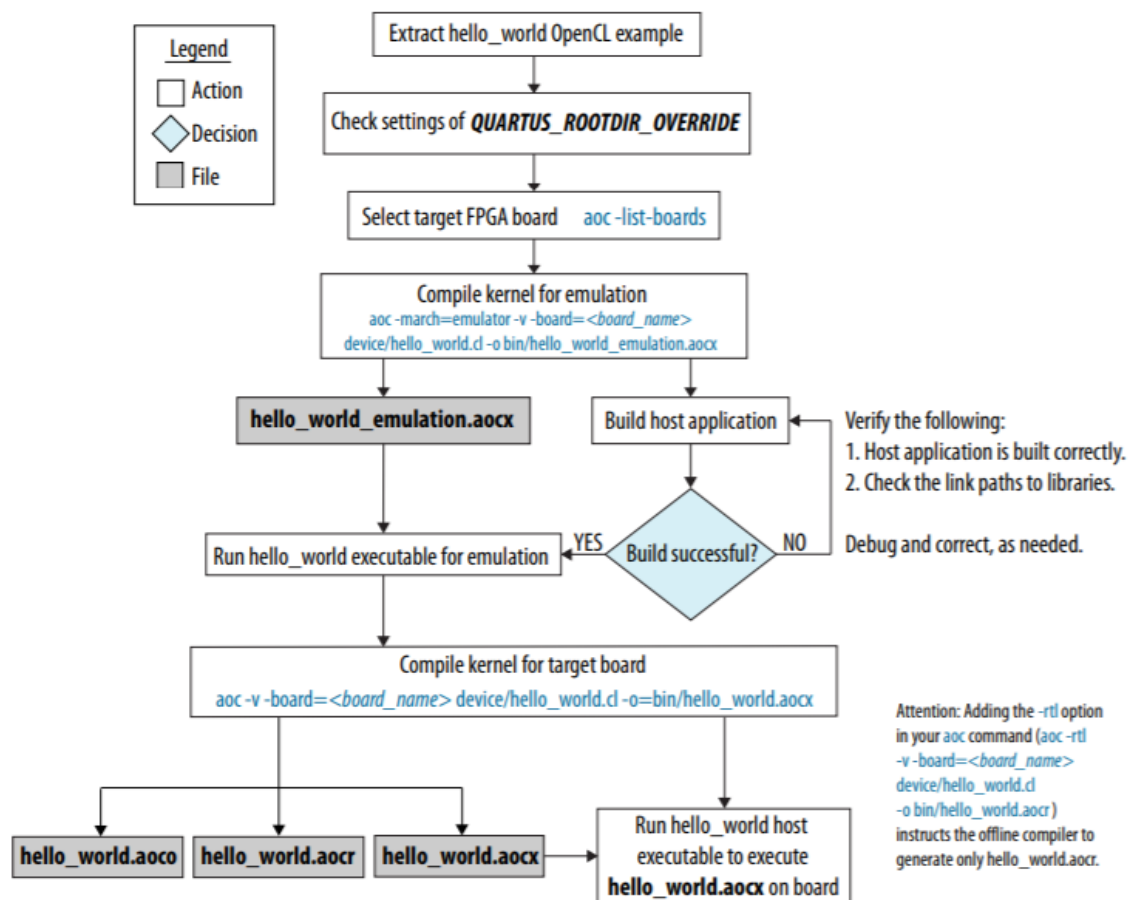
Conceptual OpenCL device architecture

- Host is not shown



The Intel FPGA OpenCL flow

- You need a Board Support Package


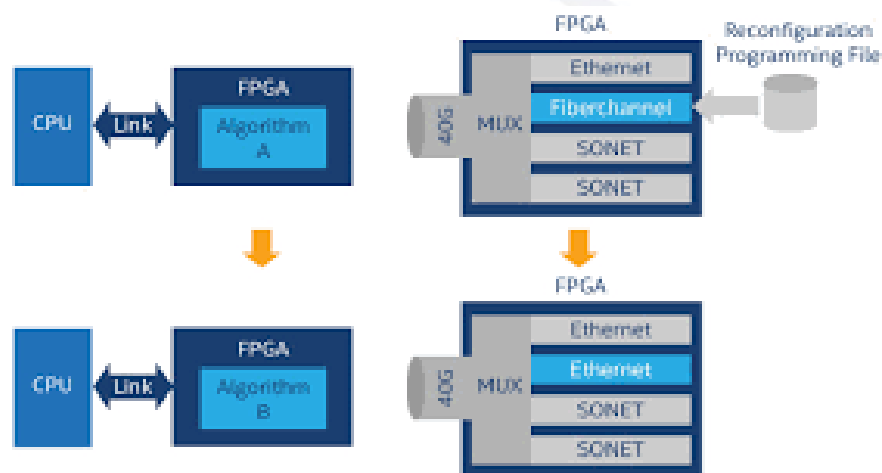


The Intel FPGA OpenCL Flow

- You have **MUCH MORE** freedom when you design hardware
- OpenCL
 - Allows fast Design Space Exploration
 - Easier to code than HDL (Verilog, VHDL)
- FPGA compilation takes long time
 - No runtime compilation (unlike GPUs)
 - Ahead of time compiler (clCreateProgramWithBinary)

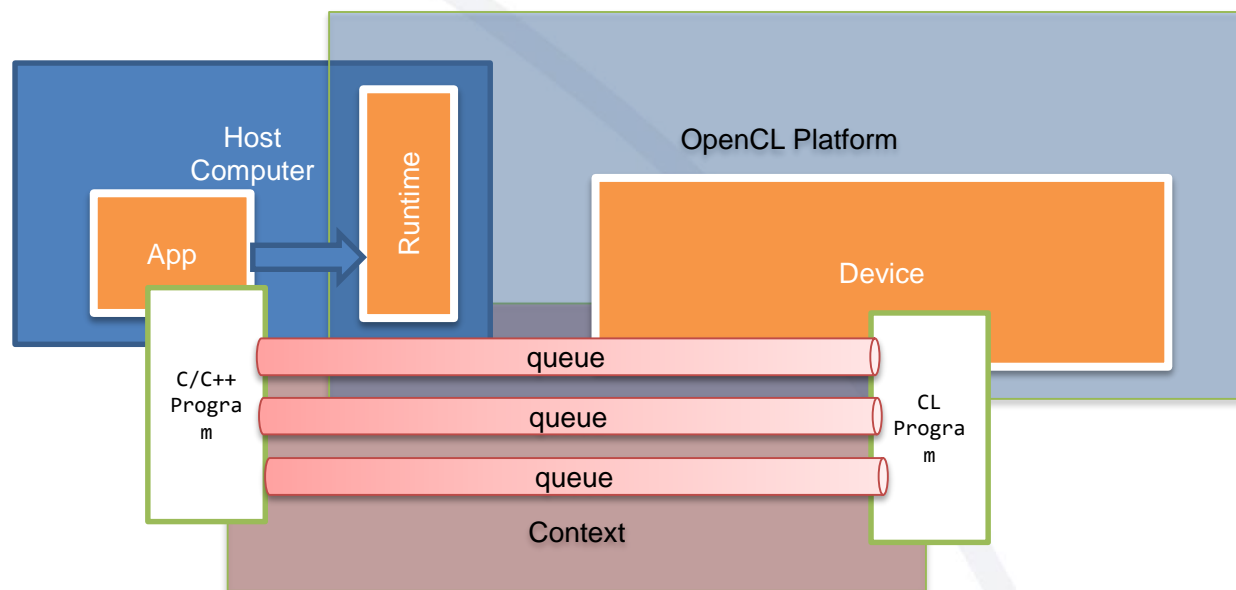
The Intel FPGA OpenCL Flow

- Partial Reconfiguration
 - Reserve a portion of the FPGA for future use
 - Change it during runtime



The Host API

- You start from C code
- There is a API
 - some headers files (CL/cl.h)
 - A dynamic library (libOpenCL.so in Linux, OpenCL.dll in Windows)
- There is a user/provider architecture
 - Several providers can implement the OpenCL API.
 - Those provide something call ICD (Installable Client Driver)
 - Intel has an ICD for (Intel OpenCL SDK)
 - NVIDIA has an ICD for (CUDA based OpenCL)
 - Cygwin has some POCL
 - In Linux & Cygwin the providers are listed in /etc/OpenCL/vendors



- The Platform Concept
 - It is a “host” system connected to a number of devices
 - It can be used to export different methods on the same physical machine
- Relevant functions
 - `clGetPlatformIDs`
 - `clGetPlatformInfo`

Example

```
void OpenCL_Interface::selectPlatform(cl_int m_platform_id)
{
    cl_int status;
    cl_uint num_platforms;
    int id_platform = m_platform_id;

    /* Get Platform Info */
    status = clGetPlatformIDs(0, nullptr, &num_platforms);
    checkError(status, "Error calling clGetPlatformIDs");
    printf("\n - Platforms (%d):", num_platforms);

    cl_platform_id platforms[num_platforms];

    status = clGetPlatformIDs(num_platforms, platforms, nullptr);
    checkError(status, "Error calling clGetPlatformIDs");

    for (cl_uint i = 0; i < num_platforms; ++i)
    {
        // Get the length for the i-th platform name
        size_t platform_name_length = 0;
        status = clGetPlatformInfo(platforms[i], CL_PLATFORM_NAME, 0,
        nullptr, &platform_name_length);
        checkError(status, "Error calling clGetPlatformInfo");

        // Get the name itself for the i-th platform
        // use vector for automatic memory management
        char platform_name[platform_name_length];
        status = clGetPlatformInfo(platforms[i], CL_PLATFORM_NAME,
        platform_name_length, platform_name, nullptr);
        checkError(status, "Error calling clGetPlatformInfo");

        printf("\n\r    [%d] %s", i, platform_name);
```

```
        if (id_platform == i)
        {
            printf(" [Selected]");

            if (strstr(platform_name, "FPGA") != NULL)
                m_isFpga = true;
        }

        fflush(stdout);
    }
    printf("\n");
    m_platform = platforms[id_platform];

    char char_buffer[STRING_BUFFER_LEN];
    printf("Detected OpenCL platforms: %d\n", num_platforms);
    printf("Querying platform for info:\n");
    printf("%42s\n", "=====");
    clGetPlatformInfo(m_platform, CL_PLATFORM_NAME,
    STRING_BUFFER_LEN, char_buffer, NULL);
    printf("%-40s = %s\n", "CL_PLATFORM_NAME", char_buffer);
    clGetPlatformInfo(m_platform, CL_PLATFORM_VENDOR,
    STRING_BUFFER_LEN, char_buffer, NULL);
    printf("%-40s = %s\n", "CL_PLATFORM_VENDOR ", char_buffer);
    clGetPlatformInfo(m_platform, CL_PLATFORM_VERSION,
    STRING_BUFFER_LEN, char_buffer, NULL);
    printf("%-40s = %s\n", "CL_PLATFORM_VERSION ", char_buffer);

    printf("\n");
}
```

- The Device Concept
 - A particular device inside the platform

Example

```
void OpenCL_Interface::selectDevice()
{
    cl_int status;
    int id_device = 0;
    cl_uint num_devices;

    /* Get Device Info */
    status = clGetDeviceIDs(m_platform, CL_DEVICE_TYPE_ALL, 0, nullptr, &num_devices);
    checkError(status, "Error calling clGetDeviceIDs");
    printf(" - Devices (%d):", num_devices);

    std::vector<cl_device_id> devices(num_devices);

    status = clGetDeviceIDs(m_platform, CL_DEVICE_TYPE_ALL, num_devices, &devices[0], nullptr);
    checkError(status, "Error calling clGetDeviceIDs");

    for (cl_uint i = 0; i < num_devices; ++i)
    {
        // Get the length for the i-th device name
        size_t device_name_length = 0;
        status = clGetDeviceInfo(devices[i], CL_DEVICE_NAME, 0, nullptr, &device_name_length);
        checkError(status, "Error calling clGetDeviceInfo");

        // Get the name itself for the i-th device
        // use vector for automatic memory management
        char device_name[device_name_length];
        status = clGetDeviceInfo(devices[i], CL_DEVICE_NAME, device_name_length, device_name, nullptr);
        checkError(status, "Error calling clGetDeviceInfo");

        printf("\n\t [%d] %s", i, device_name);

        if (id_device == i)
            printf(" [Selected]");

        fflush(stdout);
    }
    printf("\n");
    m_device = devices[id_device];
}
```


- It groups the objects to talk to a device (queues, memories, etc.)

Example

```
m_context = clCreateContext(NULL, 1, &m_device, NULL, NULL, &status);  
checkError(status, "Failed to create Context");
```

Program

- This is the program that will be executed in the Device
- You have 2 options to create a program
 - From Source Code

`cl_program` **`clCreateProgramWithSource`** (`cl_context` context,
cl_uint count, const char **strings,
const size_t *lengths, cl_int *errcode_ret)

- Standard method in GPUs
- We usually store the program in .cl files

– From Binary

- **`clCreateProgramWithBinary`**
- Standard method in FPGAs (because compilation is soooo slow)

– Programs must be compiled

Creating a program from Source

```
void OpenCL_Interface::createProgramFromSource()
{
    FILE *fp;
    char *source_str;
    size_t source_size;
    cl_int status;

    // Load the source code containing the kernel
    fp = fopen(m_kernel_file.c_str(), "r");
    if (!fp)
    {
        std::cerr << "Failed to load kernel!" << m_kernel_file.c_str() << std::endl;
        exit(EXIT_FAILURE);
    }

    source_str = (char*) malloc(MAX_SOURCE_SIZE);
    source_size = fread(source_str, 1, MAX_SOURCE_SIZE, fp);
    fclose(fp);

    // Create Kernel Program from the source
    m_program = clCreateProgramWithSource(gContext, 1, (const char **) &source_str, (const size_t *) &source_size,
    &status);
    checkError(status, "Failed to create Program with Source");

    free(source_str); //malloc of source code
}
```

Kernels

- A kernel is a function of the Device program which is accessible from the host

- identified by `__kernel` word

- An example syntax

```
__kernel void add_kernel(__global int* in_array_a, __global int* in_array_b,
                        __global int* in_array_r, int n)
{
    for (int i=0; i < n; i++)
        r[i] = a[i]+b[i];
}
```

- The Host program need to get the kernel identifiers with

```
cl_kernel clCreateKernel (cl_program program,
                        const char *kernel_name,
                        cl_int *errcode_ret);
```

Example of creating program a

```
....

char* options = "";

status = clBuildProgram(gProgram, 1, &gDevice, options, NULL, NULL);

if (status != CL_SUCCESS)
{
    size_t len;
    char buffer[4096];
    clGetProgramBuildInfo(gProgram, gDevice, CL_PROGRAM_BUILD_LOG, sizeof(buffer), buffer, &len);
    std::cout << buffer << std::endl;
    fflush(stdout);
    checkError(status, "Failed to Build Program");
}

/* Create OpenCL Kernel/s */
string kernel_name = "add_kernel";

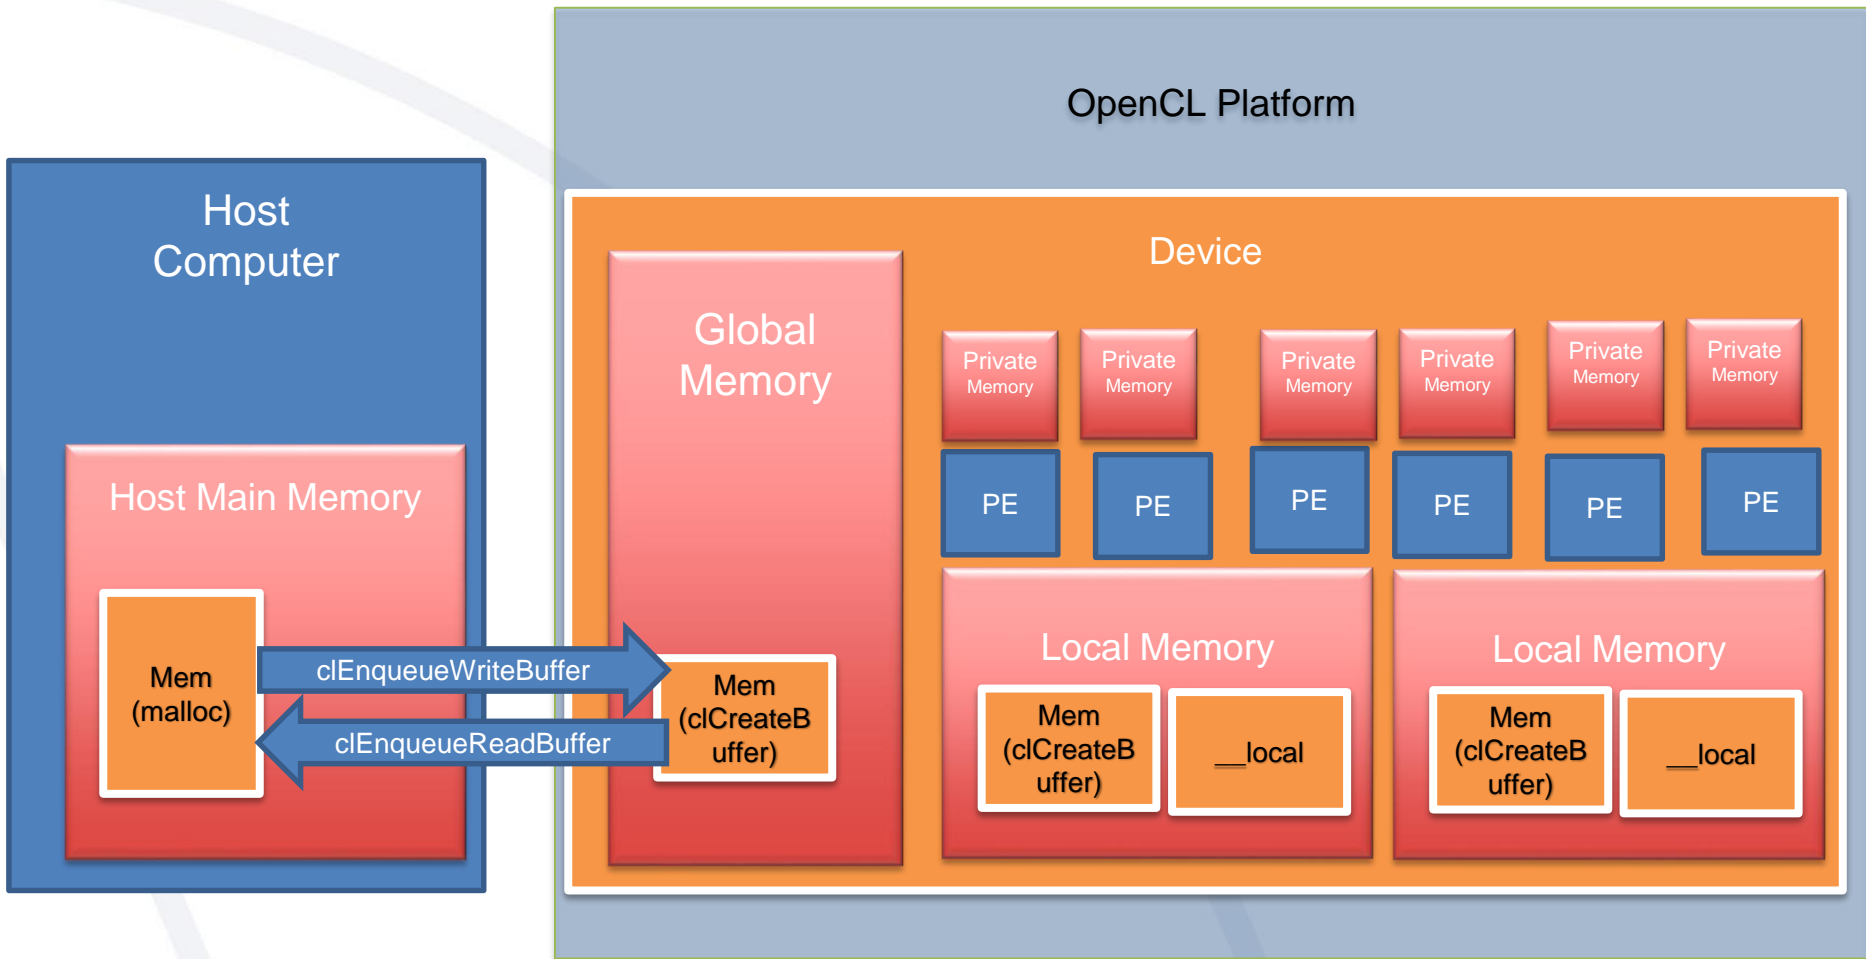
gKernel = clCreateKernel(gProgram, m_kernel_name.c_str(), &status);

char msg[100];
sprintf(msg, "Failed to Create Kernel %s", kernel_name.c_str());
checkError(status && gKernel, msg);

printf("[ OK ]\n");
```

- a pipe to send/receive information to/from a kernel
- you can have several queues

Memory



- Creates memory into the device

```
cl_mem clCreateBuffer (    cl_context context,  
    cl_mem_flags flags,  
    size_t size,  
    void *host_ptr,  
    cl_int *errcode_ret)
```

Transferring Memory

- Transfer byte from host to device

```
cl_int clEnqueueWriteBuffer (    cl_command_queue command_queue,
                                cl_mem buffer,
                                cl_bool blocking_write,  size_t offset,    size_t cb,    const void
*ptr,
                                cl_uint num_events_in_wait_list, const cl_event *event_wait_list,
                                cl_event *event)
```

```
cl_int clEnqueueReadBuffer (    cl_command_queue command_queue,
                                cl_mem buffer,
                                cl_bool blocking_read,   size_t offset,    size_t cb,    void *ptr,
                                cl_uint num_events_in_wait_list, const cl_event *event_wait_list,
                                cl_event *event)
```

Kernel Invocation

- Argument Setting

- Invoking kernel execution

```
cl_int clEnqueueNDRangeKernel (    cl_command_queue  
command_queue,  
    cl_kernel kernel, cl_uint work_dim,  
    const size_t *global_work_offset,    const size_t  
*global_work_size,  
    const size_t *local_work_size, cl_uint  
num_events_in_wait_list,  
    const cl_event *event_wait_list,    cl_event  
*event)
```

Work Items and Workgroups

- The work units in OpenCL are called Work-items
 - They are equivalent to CUDA threads.
- Work-items can be grouped in Workgroups.
 - Workgroups share local memory.
 - They can be synchronized by barriers
- Work items are identified by IDs.
 - `global_id` is unique at the global scale
 - `local_id` is unique at the workgroup scale
- The workgroup size is defined explicitly with the local size, but the number of workgroups is inferred from global and local dimensions when calling `clEnqueueNDRangeKernel`
 - $\text{Num Workgroups} = \text{GlobalSize} / \text{LocalSize}$

Wort Items and Workgroups

- `get_global_id(dim)` (`gid`)
 - returns the the global ID in the `dim` dimension (we can have up to 3 dimensions)
- `get_local_id(dim)` (`lid`)
 - returns the the local ID in the `dim` dimension (we can have up to 3 dimensions)
- `get_group_id(dim)` (`wid`)
 - returns the the work-group ID in the `dim` dimension (we can have up to 3 dimensions)

$$\text{gid} = \text{wid} * \text{wsz} + \text{lid}$$

- The workgroup size is not defined explicitly but inferred by global and local dimensions when calling `clEnqueueNDRangeKernel`
 - Number of Workgroups = $\text{GlobalSize} / \text{LocalSize}$


The Kernel API

Kernel Code

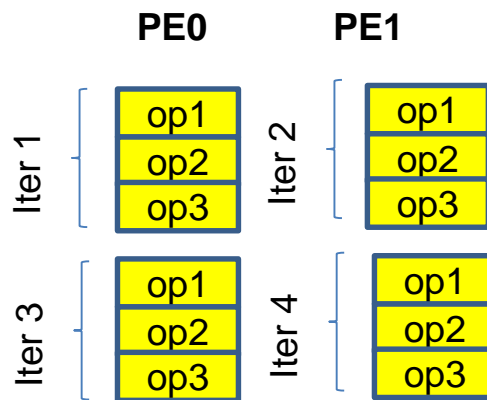
- Just regular C
- Functions
 - To be a kernel you just add `__kernel` before function declaration
 - Kernels functions are void (no return parameter)
- NO STACK -> ALWAYS INLINING
- Tricks are played with kernel APIs (`get_global_id`, `barriers`, etc), `#pragma` and more...

Loop Unrolling and Loop pipelining

```
for each (iter ){
    op1;
    op2;
    op3;
}
```



• Unrolling



• Pipelining

